

# ensemble learning

October 24, 2022

```
[1]: import numpy as np
import pandas as pd
```

```
[2]: data= pd.read_csv("horse.csv")
```

```
[3]: data.head()
```

```
[3]:  surgery    age  hospital_number  rectal_temp  pulse  respiratory_rate  \
0      no  adult         530101         38.5    66.0             28.0
1     yes  adult         534817         39.2    88.0             20.0
2      no  adult         530334         38.3    40.0             24.0
3     yes  young         5290409         39.1   164.0             84.0
4      no  adult         530255         37.3   104.0             35.0

    temp_of_extremities  peripheral_pulse  mucous_membrane  capillary_refill_time  \
0                cool          reduced              NaN      more_3_sec
1                NaN              NaN  pale_cyanotic      less_3_sec
2              normal          normal      pale_pink      less_3_sec
3                cold          normal  dark_cyanotic      more_3_sec
4                NaN              NaN  dark_cyanotic      more_3_sec

    ...  packed_cell_volume  total_protein  abdomo_appearance  abdomo_protein  \
0  ...              45.0           8.4              NaN              NaN
1  ...              50.0          85.0             cloudy              2.0
2  ...              33.0           6.7              NaN              NaN
3  ...              48.0           7.2      serosanguinous              5.3
4  ...              74.0           7.4              NaN              NaN

    outcome  surgical_lesion  lesion_1  lesion_2  lesion_3  cp_data
0      died              no     11300         0         0       no
1  euthanized              no      2208         0         0       no
2      lived              no         0         0         0      yes
3      died              yes      2208         0         0      yes
4      died              no      4300         0         0       no
```

```
[5 rows x 28 columns]
```

```
[4]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 299 entries, 0 to 298
Data columns (total 28 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   surgery                              299 non-null    object
1   age                                  299 non-null    object
2   hospital_number                      299 non-null    int64
3   rectal_temp                          239 non-null    float64
4   pulse                                275 non-null    float64
5   respiratory_rate                     241 non-null    float64
6   temp_of_extremities                  243 non-null    object
7   peripheral_pulse                     230 non-null    object
8   mucous_membrane                      252 non-null    object
9   capillary_refill_time                267 non-null    object
10  pain                                  244 non-null    object
11  peristalsis                          255 non-null    object
12  abdominal_distention                 243 non-null    object
13  nasogastric_tube                     195 non-null    object
14  nasogastric_reflux                   193 non-null    object
15  nasogastric_reflux_ph                53 non-null     float64
16  rectal_exam_feces                   197 non-null    object
17  abdomen                              181 non-null    object
18  packed_cell_volume                   270 non-null    float64
19  total_protein                        266 non-null    float64
20  abdomo_appearance                    134 non-null    object
21  abdomo_protein                       101 non-null    float64
22  outcome                              299 non-null    object
23  surgical_lesion                      299 non-null    object
24  lesion_1                             299 non-null    int64
25  lesion_2                             299 non-null    int64
26  lesion_3                             299 non-null    int64
27  cp_data                              299 non-null    object
dtypes: float64(7), int64(4), object(17)
memory usage: 65.5+ KB
```

```
[5]: data.shape
```

```
[5]: (299, 28)
```

```
[6]: data.size
```

```
[6]: 8372
```

```
[7]: data.isna().sum()
```

```

[7]: surgery          0
    age              0
    hospital_number   0
    rectal_temp       60
    pulse             24
    respiratory_rate   58
    temp_of_extremities 56
    peripheral_pulse    69
    mucous_membrane    47
    capillary_refill_time 32
    pain              55
    peristalsis        44
    abdominal_distention 56
    nasogastric_tube   104
    nasogastric_reflux 106
    nasogastric_reflux_ph 246
    rectal_exam_feces 102
    abdomen           118
    packed_cell_volume 29
    total_protein      33
    abdomo_appearance 165
    abdomo_protein     198
    outcome            0
    surgical_lesion     0
    lesion_1           0
    lesion_2           0
    lesion_3           0
    cp_data            0
    dtype: int64

```

```
[9]: features= data.drop("outcome" , axis=1)
```

```
[10]: target= data['outcome']
```

```
[11]: features_transformed= pd.get_dummies(features)
```

```
[12]: from sklearn.model_selection import train_test_split
```

```
[13]: from sklearn.tree import DecisionTreeClassifier
```

```
[14]: from sklearn.ensemble import RandomForestClassifier
```

```
[15]: x_train, x_test, y_train, y_test = train_test_split(features_transformed,
    ↪target, random_state=2)
```

```
[16]: print(features_transformed.shape, x_train.shape, x_test.shape)
```

(299, 67) (224, 67) (75, 67)

```
[18]: from sklearn.impute import SimpleImputer
```

```
[20]: imputer = SimpleImputer(missing_values= np.nan, strategy = 'most_frequent')
```

```
[21]: x_train= imputer.fit_transform(x_train)
      x_test = imputer.fit_transform(x_test)
```

```
[37]: my_tree = DecisionTreeClassifier(criterion = 'entropy', random_state= 2)
```

```
[38]: my_tree.fit(x_train, y_train)
```

```
[38]: DecisionTreeClassifier(criterion='entropy', random_state=2)
```

```
[39]: ypred= my_tree.predict(x_test)
```

```
[40]: from sklearn.metrics import accuracy_score, confusion_matrix, \
      ↪classification_report
```

```
[41]: accuracy_score(y_test, ypred)
```

```
[41]: 0.56
```

```
[42]: confusion_matrix(y_test, ypred)
```

```
[42]: array([[ 8,  2, 16],
          [ 1,  2,  5],
          [ 3,  6, 32]])
```

```
[44]: print(classification_report(y_test, ypred))
```

	precision	recall	f1-score	support
died	0.67	0.31	0.42	26
euthanized	0.20	0.25	0.22	8
lived	0.60	0.78	0.68	41
accuracy			0.56	75
macro avg	0.49	0.45	0.44	75
weighted avg	0.58	0.56	0.54	75

### voting classifier

```
[47]: from sklearn.ensemble import VotingClassifier
      from sklearn.linear_model import LogisticRegression
      from sklearn.svm import SVC
```

```
[48]: lr_class= LogisticRegression()  
      svc_class= SVC()  
      rf_class= RandomForestClassifier()
```

```
[50]: voting_classifier= VotingClassifier(estimators =[('lr',lr_class ), ('svc',  
→svc_class),('rf',rf_class)])
```

```
[51]: voting_classifier.fit(x_train,y_train)
```

```
/usr/local/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

```
[51]: VotingClassifier(estimators=[('lr', LogisticRegression()), ('svc', SVC()),  
                                  ('rf', RandomForestClassifier())])
```

```
[57]: for clf in (lr_class,svc_class,rf_class,voting_classifier):  
      clf.fit(x_train,y_train)  
      y_pred = clf.predict(x_test)  
      print(clf.__class__.__name__, accuracy_score(y_test, y_pred))
```

```
/usr/local/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

LogisticRegression 0.5466666666666666

SVC 0.5466666666666666

RandomForestClassifier 0.68

VotingClassifier 0.5466666666666666

```
/usr/local/lib/python3.7/site-packages/sklearn/linear_model/_logistic.py:818:  
ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
 Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
 regression  
 extra\_warning\_msg=\_LOGISTIC\_SOLVER\_CONVERGENCE\_MSG,

## 0.1 Bagging and Boosting

```
[58]: from sklearn.ensemble import BaggingClassifier
```

```
[60]: bag_clf= BaggingClassifier(DecisionTreeClassifier(), n_estimators= 100)
```

```
[61]: bag_clf.fit(x_train,y_train)
```

```
[61]: BaggingClassifier(base_estimator=DecisionTreeClassifier(), n_estimators=100)
```

```
[63]: y_pred= bag_clf.predict(x_test)
```

```
[65]: accuracy_score(y_test, y_pred)
```

```
[65]: 0.6133333333333333
```

```
[66]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
died	0.61	0.65	0.63	26
euthanized	0.00	0.00	0.00	8
lived	0.67	0.71	0.69	41
accuracy			0.61	75
macro avg	0.43	0.45	0.44	75
weighted avg	0.58	0.61	0.60	75

```
[67]: confusion_matrix(y_test, y_pred)
```

```
[67]: array([[17,  0,  9],
        [ 3,  0,  5],
        [ 8,  4, 29]])
```

```
[68]: from sklearn.ensemble import AdaBoostClassifier
```

```
[69]: ada_boost=AdaBoostClassifier(DecisionTreeClassifier(), n_estimators=100)
```

```
[70]: ada_boost.fit(x_train,y_train)
```

```
[70]: AdaBoostClassifier(base_estimator=DecisionTreeClassifier(), n_estimators=100)
```

```
[71]: y_pred= ada_boost.predict(x_test)
```

```
[72]: accuracy_score(y_test, y_pred)
```

```
[72]: 0.5466666666666666
```

```
[73]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
died	0.52	0.58	0.55	26
euthanized	0.12	0.12	0.12	8
lived	0.66	0.61	0.63	41
accuracy			0.55	75
macro avg	0.43	0.44	0.43	75
weighted avg	0.55	0.55	0.55	75

```
[74]: confusion_matrix(y_test, y_pred)
```

```
[74]: array([[15,  2,  9],  
        [ 3,  1,  4],  
        [11,  5, 25]])
```

```
[75]: from sklearn.ensemble import GradientBoostingClassifier
```

```
[77]: grad_boost= GradientBoostingClassifier()
```

```
[78]: grad_boost.fit(x_train,y_train)
```

```
[78]: GradientBoostingClassifier()
```

```
[79]: y_pred=grad_boost.predict(x_test)
```

```
[80]: accuracy_score(y_test, y_pred)
```

```
[80]: 0.56
```

```
[81]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
died	0.53	0.69	0.60	26

euthanized	0.00	0.00	0.00	8
lived	0.69	0.59	0.63	41
accuracy			0.56	75
macro avg	0.41	0.43	0.41	75
weighted avg	0.56	0.56	0.55	75

```
[82]: confusion_matrix(y_test, y_pred)
```

```
[82]: array([[18,  1,  7],
           [ 4,  0,  4],
           [12,  5, 24]])
```

```
[87]: import xgboost
```

```
[88]: xgb_class= xgboost.XGBClassifier()
```

```
[89]: xgb_class.fit(x_train,y_train)
```

```
[89]: XGBClassifier(base_score=0.5, booster=None, colsample_bylevel=1,
                    colsample_bynode=1, colsample_bytree=1, gamma=0, gpu_id=-1,
                    importance_type='gain', interaction_constraints=None,
                    learning_rate=0.300000012, max_delta_step=0, max_depth=6,
                    min_child_weight=1, missing=nan, monotone_constraints=None,
                    n_estimators=100, n_jobs=0, num_parallel_tree=1,
                    objective='multi:softprob', random_state=0, reg_alpha=0,
                    reg_lambda=1, scale_pos_weight=None, subsample=1,
                    tree_method=None, validate_parameters=False, verbosity=None)
```

```
[90]: y_pred=xgb_class.predict(x_test)
```

```
[91]: accuracy_score(y_test, y_pred)
```

```
[91]: 0.6266666666666667
```

```
[92]: confusion_matrix(y_test, y_pred)
```

```
[92]: array([[17,  1,  8],
           [ 3,  1,  4],
           [ 9,  3, 29]])
```

```
[93]: print(classification_report(y_test, y_pred))
```

	precision	recall	f1-score	support
died	0.59	0.65	0.62	26
euthanized	0.20	0.12	0.15	8



lived	0.71	0.71	0.71	41
accuracy			0.63	75
macro avg	0.50	0.50	0.49	75
weighted avg	0.61	0.63	0.62	75

```
[ ]:
```